

FULL-STACK AI ENGINEER

LUIS DIEGO COLÍN MENDIOLA

Product-Driven | Agentic Systems | Production AI

 luiscolin764@gmail.com

 [LinkedIn](#)

 [Portfolio](#)

 +52 715 145 9328

ABOUT ME

Product-focused Full-Stack AI Engineer specializing in **Agentic Workflows (LangGraph)** and **High-Performance Backends (FastAPI)**. Winner of the **Meta Llama Impact Grant (\$100k)** with experience shipping production RAG, fine-tuning, and edge AI systems for startups and enterprise.

SPECIALIZED TRAINING

- Meta Llama Ecosystem For AI Development | Meta, CENTRO, Infotech
- LLM Engineering: Large Language Models & Autonomous Agents - *Advanced Certification* | 2025

HARD SKILLS

- **AI & ML:** Agentic AI (LangGraph, LangChain), RAG Pipelines, LLM Fine-tuning (LoRA), Edge LLM Deployment (PyTorch ExecuTorch / PTE), RunPod, OpenAI API
- **Backend:** Python, FastAPI, Pydantic, PostgreSQL (SQL), Redis, Celery
- **Frontend:** Nextjs, React Native, Tailwind CSS, TypeScript
- **Languages:** English (C1), Spanish (Native)
- **Deployment:** Vercel, Railway, Docker, Git, GitHub, CI/CD Workflows

EXPERIENCE

Co-Founder & Lead AI Engineer | BluAI | 2024 - Present

Winner – Meta Llama Impact Grant (\$100k USD)

- Co-founded and lead technical development for BluAI, a hurricane resilience platform launching to production on Google Play in May 2026, with iOS in active development.
- Architected a Hybrid AI System: Llama 3.3 (Cloud) with RAG and tool-calling for complex reasoning, switching to Quantized Llama Edge Models (PTE) for fully offline inference during connectivity blackouts.
- Engineered the full-stack mobile app in React Native + FastAPI, integrating map-based risk visualization, personalized disaster alerts, and a hybrid cloud/edge inference pipeline.

Full-Stack AI Engineer | Dataframe AI Spine | 2025

- Engineered the core architecture for AI SPINE, a platform for building autonomous agents using natural language.
- Implemented complex agentic memory and tool-calling workflows using LangChain and LangGraph, enabling agents to retain context across sessions and platforms.
- Delivered a modular, production-ready MVP (FastAPI/Python) and executed a full technical handover to the founding team. Retained as technical advisor.

Full-Stack AI Engineer | Harris & Frank | 2025

- Built a production web app for sales teams that generates photorealistic suit previews from raw fabric textures, replacing physical photoshoots.
- Fine-tuned custom SDXL + LoRA models on RunPod; full-stack architecture in Next.js + FastAPI + Cloudflare R2.

Personal Project | The Overseer | 2026 - Present

- Built a personal LangGraph agent that enforces task deadlines by polling a source of truth (Notion API), then escalates interventions via Twilio (WhatsApp → calls) as deadlines approach.
- Engineered the backend with FastAPI, PostgreSQL (SQLAlchemy), and APScheduler for persistent task state and scheduled polling; serves as a personal tool.

Co-Founder | Verskod | 2024 - 2026

- Co-founded a development consultancy delivering full-stack and AI engineering for early-stage startups. Led technical scoping and **Web + FastAPI + PostgreSQL** builds. Exited as co-founder; retained as technical advisor.